

# Concurrent and Parallel Programming, Part II

Programming Languages

CS 214




# Distributed Synchronization

Semaphores, locks, condition variables, monitors, are *shared-memory* constructs, and so \_\_\_\_\_.

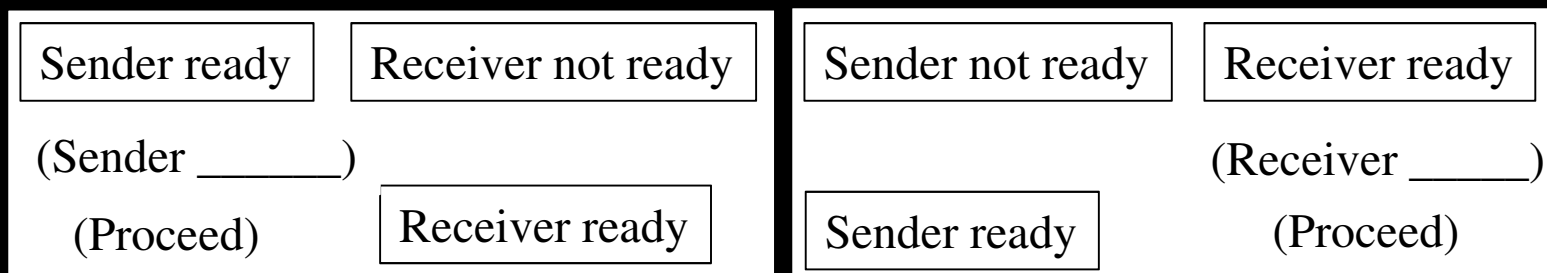
- They are of \_\_\_\_\_ on a *distributed multiprocessor*

On a distributed multiprocessor, processes can communicate via \_\_\_\_\_ primitives.

- If the message-passing system has *no storage*, then the send/receive operations must be *synchronized*:




Two scenarios...



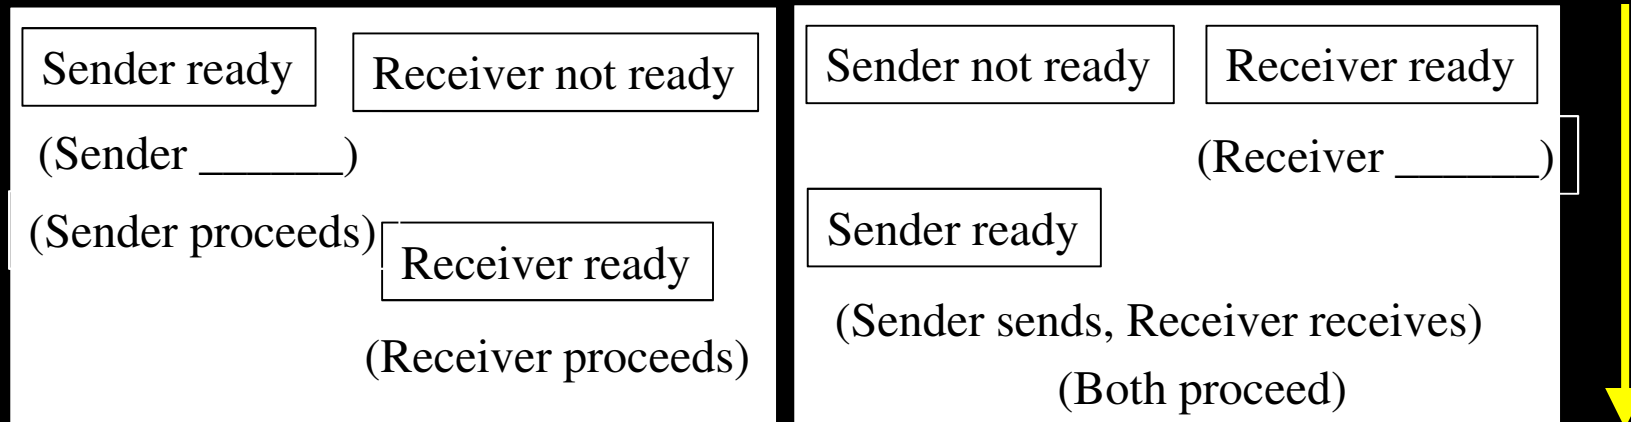
# Asynchronous Communication

–If the message-passing system has \_\_\_\_\_,  
then the send/receive operations can proceed \_\_\_\_\_:



The receiver can then retrieve the message when it is ready...

Two scenarios...



Message-buffering eliminates some (but not all) of the waiting.



# Message-Passing Languages

Some languages support message-passing between \_\_\_\_\_:

- \_\_\_\_\_ is a functional language developed at Ericcson and used by Nortel, T-Mobile, Facebook (chat, WhatsApp), and 20+ others.



```
B!{self(), {digits, [2,3,5]}}
```

```
receive  
  {A, {digits, nums}} ->  
    analyze(nums);  
end
```

- \_\_\_\_\_ is a hybrid OO+functional language used at Netflix, LinkedIn, Twitter, Tumblr, Foursquare, Sony, and other companies:

```
B! digits(2,3,5)
```

```
receive {  
  case digits(nums) =>  
    analyze(nums);  
}
```



# An Ada Task

... has 3 characteristics:

- its own \_\_\_\_\_ of control;
- its own \_\_\_\_\_; and
- \_\_\_\_\_ (aka *entry procedures*)

Entry procedures are *self-synchronizing subprograms* that another task can invoke for task-to-task communication.

If task *t1* has an entry procedure *p*, then another task *t2* can execute:

*t1.p( argument-list );*

In order for *p* to execute, *t1* must execute:

*accept p ( parameter-list );*

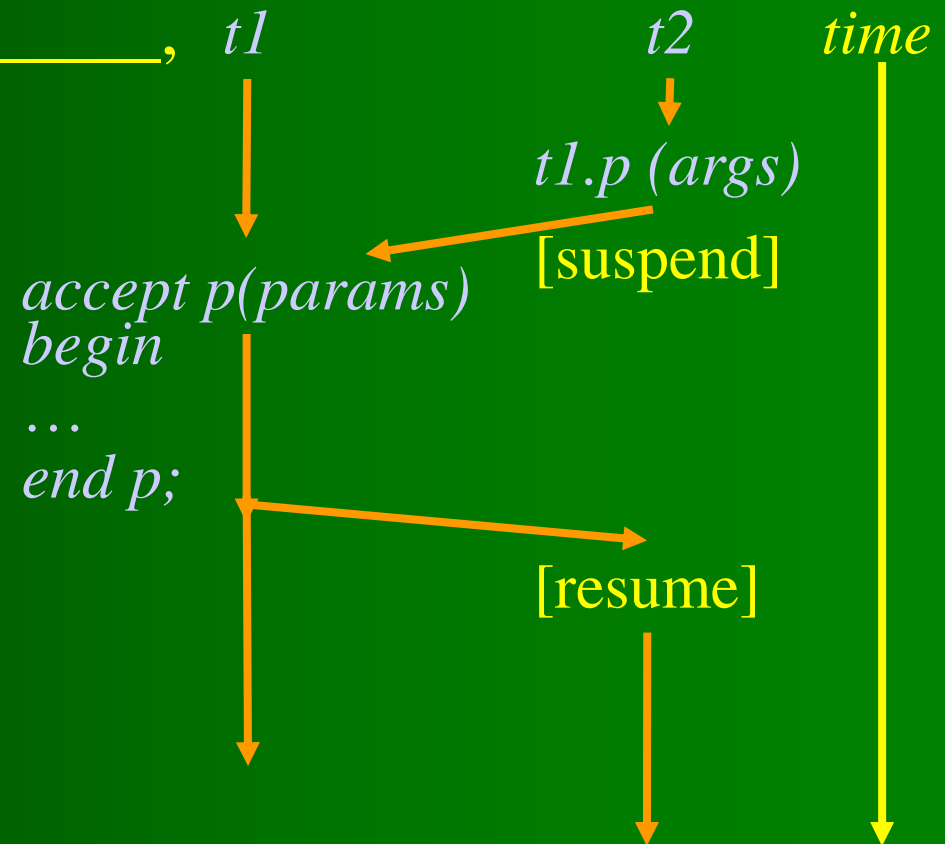
- If *t1* executes *accept p* and *t2* has not called *p*, \_\_\_\_\_;
- If *t2* calls *p* and *t1* has not done *accept p*, \_\_\_\_\_.



# Rendezvous

When  $t1$  and  $t2$  are \_\_\_\_\_,  $t1$  and  $t2$  execute:

- $t2$ 's *argument-list* is evaluated and passed to  $t1.p$ 's parameters
- $t2$  suspends
- $t1$  executes the body of  $p$ , using its parameter values
- return-values (or *out* or *in out* parameters) are passed back to  $t2$
- $t1$  continues execution;  $t2$  resumes execution



This interaction is called a \_\_\_\_\_ between  $t1$  and  $t2$ .  
It does not depend on shared memory, so  $t1$  and  $t2$  can be on a uniprocessor, a tightly-coupled or a distributed multiprocessor.



# Ada Array Processing

How can we rewrite what's below to complete more quickly?

```
procedure sumArray is
  N: constant integer := 1000000;
  type RealArray is array(1..N) of float;
  anArray: RealArray;

  function sum(a: RealArray; first, last: integer)
    return float is
    result: float := 0.0;
  begin
    for i in first..last loop
      result := result + a(i);
    end loop;
    return result;
  end sum;

begin
  -- code to fill anArray with values omitted
  put( sum(anArray, 1, N) );
end sumArray;
```



# Divide-And-Conquer via Tasks

```
procedure parallelSumArray is
  -- declarations of N, RealArray, anArray, Sum() as before ...
```

```
  task type ArraySliceAdder
    entry SumSlice(Start: in Integer; Stop: in Integer);
    entry GetSum(Result: out float);
  end ArraySliceAdder;
```

```
  task body ArraySliceAdder is
    i, j: Integer; Answer: Float;
  begin
```

```
    accept SumSlice(Start: in Integer; Stop: in Integer) do
      i := Start; j := Stop;           -- get inputs
    end SumSlice;
```

```
    Answer := Sum(anArray, i, j);      -- do the work
```

```
    accept GetSum(Result: out float) do
      Result := Answer;                -- report outcome
    end GetSum;
```

```
  end ArraySliceAdder;
```

```
-- continued on next slide...
```





# Divide-And-Conquer via Tasks (ii)

```
-- continued from previous slide ...

firstHalfSum, secondHalfSum: Integer;
T1, T2 : ArraySliceAdder;    -- T1, T2 start & wait on accept
begin
  -- code to fill anArray with values omitted

  T1.SumSlice(1, N/2);        -- start T1 on 1st half
  T2.SumSlice(N/2 + 1, N);    -- start T2 on 2nd half

  T1.GetSum( firstHalfSum );  -- get 1st half sum from T1
  T2.GetSum( secondHalfSum ); -- get 2nd half sum from T2

  put( firstHalfSum + secondHalfSum ); -- we're done!
end parallelSumArray;
```

Using two tasks T1 and T2, this *parallelSumArray* version requires roughly 1/2 the time required by *sumArray* (on a multiprocessor).

Using three tasks, the time will be roughly 1/3 the time of *sumArray*.

...



# Producer-Consumer in Ada

To give the producer and consumer separate threads, we can define the behavior of one in the 'main' procedure:

and the behavior of the other in a separate task:

We can then build a Monitor-style *Buffer task* with *put()* and *get()* as (auto-synchronizing) entry procedures...

```
procedure ProducerConsumer is
  buf: Buffer;
  it: Item;
```

```
task consumer;
task body consumer is
  it: Item;
begin
  loop
    buf.get(it);
    -- consume Item it
  end loop;
end consumer;
```

```
begin -- producer task
  loop
    -- produce an Item in it
    buf.put(it);
  end loop;
end ProducerConsumer;
```



# Capacity-1 Buffer

A single-value buffer is easy to build using an Ada \_\_\_\_\_:

As a *task-type*, variables of this type (e.g., *buf*) will automatically have their own thread of execution.

The body of the task is a loop that accepts calls to *put()* and *get()* in strict alternation.

This causes *buf* to alternate between being empty and nonempty.

```
task type BoundedBuffer1 is
  entry get(it: out Item);
  entry put(it: in Item);
end BoundedBuffer1;
```

```
task body BoundedBuffer1 is
  myBuffer: Item;
begin
```

```
  loop
    accept put(it: in Item) do
      myBuffer := it;
    end put;

    accept get(it: out Item) do
      it := myBuffer;
    end get;
  end loop;
```

```
end BoundedBuffer1;
```



# Capacity-N Buffer

An N-value buffer is a bit more work:

We can accept any call to *get()* so long as we are not empty, and any call to *put()* so long as we are not full.

Ada provides the *select-when* statement to \_\_\_\_\_, and \_\_\_\_\_, and perform it if and only if a given condition is *true*

```
-- task declaration is as before ...  
task body BoundedBuffer is  
  N: constant integer := 1024;  
  package Buf is new Queue(N, Items);  
begin  
  loop  
    select  
      when not Buf.isFull =>  
        accept put(it: in Item) do  
          Buf.append(it);  
        end put;  
      or when not Buf.isEmpty =>  
        accept get(it: out Item) do  
          it := Buf.first;  
          Buf.delete;  
        end get;  
      end select;  
    end loop;  
end BoundedBuffer;
```



# MPI ...

... is the \_\_\_\_\_

... is an industry-standard library for distributed-memory parallel computing in C, C++, Fortran, with 3<sup>rd</sup> party bindings for Java, Python, R, ...

... was designed by a large consortium in 1994:

- 12 companies: *Cray, IBM, Intel, ...*
- 11 national labs: *ANL, LANL, LLNL, ORNL, Sandia, ...*
- representatives from 16 universities

... has “built in” support for many parallel design patterns

... continues to evolve (MPI 2.0 in 1997; 3.0 in 2012; ...)



# Typical MPI Program Structure

```
#include <mpi.h>                                // MPI functions

int main(int argc, char** argv) {
    int id = -1, numProcesses = -1;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);

    // program body, which usually includes
    // calls to MPI_Send() and MPI_Receive()

    MPI_Finalize();
    return 0;
}
```



# The 6 MPI Basic Functions

1. `MPI_Init(&argc, &argv);`
  - Set up **MPI\_COMM\_WORLD**, a “communicator”  
(The set of processes that make up the distr. computation)
2. `MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);`
  - How many of us processes are there to attack the problem?
3. `MPI_Comm_rank(MPI_COMM_WORLD, &id);`
  - Which of the  $n$  processes am I?



# The 6 MPI Basic Functions (Ct'd)

4. `MPI_Send(sendAddress, numItems, itemType,  
destinationID, tag, communicator);`

- Send the item(s) at *sendAddress* to *destinationRank*

5. `MPI_Recv(receiveBuffer, bufferSize, itemType,  
senderID, tag, communicator, status);`

- Receive up to *bufferSize* items from *senderRank*

6. `MPI_Finalize();`

- Shut down distributed computation

These 6 commands are all you need to do useful work!





# Other Useful Functions

```
MPI_Bcast(receiveBuffer, bufferSize, itemType,  
          senderID, tag, comm, status);
```

- Broadcast *bufferSize* items from *senderID* to everyone in *comm*

```
MPI_Reduce(sendAddress, receiveBuffer, numItems,  
           itemType, combineOp, rootID,  
           tag, comm);
```


- Use *combineOp* to combine the distributed items at *sendAddress* into *receiveBuffer* at *destinationRank*

These (and many other) commands provide simple but  
efficient *collective communication*...



# Reduction (8 Processes)

To sum the local values of  $N = 8$  processes:



Reduction reduces the sum-time from  $O(N)$  to \_\_\_\_\_...



# A Very Simple MPI Program

```
#include <iostream>
#include <mpi.h>
using namespace std;

int main(int argc, char** argv) {
    int id = -1, n = -1;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &n);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);

    int startValue = id+1;
    int square = startValue * startValue;
    int sumSquares = 0;

    MPI_Reduce(&square, &sumSquares, 1, MPI_INT,
               MPI_SUM, 0, MPI_COMM_WORLD);

    if (id == 0) {
        cout << "\nThe sum of the squares from 1 to "
              << n << " is " << sumSquares << endl;
    }

    MPI_Finalize();
    return 0;
}
```



# MPI Build and Run

To build an MPI C++ *program* from the command line:

```
mpiCC program.cpp -o program
```

To run an MPI *program* from the command line:

```
mpirun -np N -machinefile hostFile ./program
```

Launch *N* processes  
(each will get a unique rank)  
Vary *N* to test scalability

Launch those *N* processes  
on the computers listed in *hostFile*  
(optional ...)

Each process runs  
this same *program*  
(SPMD pattern)



# Testing sumSquares

```
$ mpirun -np 1 ./sumSquares
```

```
The sum of the squares from 1 to 1 is 1
```

```
$ mpirun -np 2 ./sumSquares
```

```
The sum of the squares from 1 to 2 is 5
```

```
$ mpirun -np 3 ./sumSquares
```

```
The sum of the squares from 1 to 3 is 14
```

```
$ mpirun -np 4 ./sumSquares
```

```
The sum of the squares from 1 to 4 is 30
```

```
$ mpirun -np 128 ./sumSquares
```

```
The sum of the squares from 1 to 128 is 707264
```



# Summary

- *Concurrent computations* consist of multiple entities:
  - \_\_\_\_\_ in Smalltalk, MPI      – \_\_\_\_\_ in Ada
  - \_\_\_\_\_ in C/C++, C#, Java, Go, Python, Ruby, Scala, ...
- On a shared-memory multiprocessor:
  - The \_\_\_\_\_ was the first synchronization primitive
    - *Java* provides a *Semaphore* class for synchronizing processes
  - \_\_\_\_\_ and \_\_\_\_\_ separate a semaphore's mutual-exclusion usage (locks) from its synchronization usage (c.v.s)
  - \_\_\_\_\_ are higher-level, self-synchronizing objects
    - *Java* classes have an associated (simplified) monitor
- On a distributed system:
  - *Ada* tasks provide self-synchronizing \_\_\_\_\_
  - *Erlang*, *Scala*, *MPI* support \_\_\_\_\_ between processes



# Summary (ii)

## Comparing Monitors and Tasks/Threads (and Coroutines):

|             | Has Its Own Thread | Has Its Own Execution State |
|-------------|--------------------|-----------------------------|
| Monitor     |                    |                             |
| Task/Thread |                    |                             |
| Coroutine   |                    |                             |

A *coroutine* (Simula, Lua) is two or more procedures that *share a single thread*, each exercising mutual control over the other:

```
procedure A;  
begin  
  -- do something  
  resume B;  
  -- do something  
  resume B;  
  -- do something  
  -- ...  
end A;
```

```
procedure B;  
begin  
  -- do something  
  resume A;  
  -- do something  
  resume A;  
  -- ...  
end B;
```

