

Concurrent and Parallel Programming, Part I

Programming Languages

CS 214



Introduction

Suppose you and that “special someone” are shopping for champagne and oysters for a romantic dinner...

You might do your shopping either of two ways:

- The “we can’t bear to be apart” approach: together, you
 - find the oysters,
 - find the champagne,
 - pay for what you’ve selected at the checkout
- The “unromantic but efficient” divide-and-conquer approach:
 - one of you finds the oysters,
 - the other finds the champagne,
 - rendezvous at the checkout to pay for what you’ve selected

If the average time to find champagne+oysters sequentially is τ , then finding them *concurrently* takes about ____.



Concurrent Programming

Most modern programming languages provide built-in support for concurrent processing.

- _____ provides a *Task* construct, and each distinct task is automatically executed concurrently.
- _____ provides a *Thread* class, and each distinct thread can be executed concurrently.
- _____ added a standard *thread* class in C++11.
- _____, ... provide multithreading/processing capabilities

Older languages:

- _____ relies on external libraries for concurrency (e.g., Unix *fork()*, POSIX *pthreads*, OpenMP, MPI, ...).
- _____ provides a *start-process* function, but it is *not standard Lisp*.



Example: Ada

We might represent our sequential approach as follows:

```
procedure RomanticApproach is
begin
  FindOysters; ←————— do this
  FindChampagne; ←—————
  PayAtCheckout; ←—————
end RomanticApproach;
```

By contrast, our concurrent approach is:

```
procedure EfficientApproach is
  task OysterFinder;
  task body OysterFinder is begin
    FindOysters; RendezvousAtCheckout; } do this
  end OysterFinder;
begin
  FindChampagne; RendezvousAtCheckout; }
end EfficientApproach;
```



Example: Java

In Java, our sequential approach is expressed as:

```
class RomanticApproach {  
    public static void main(String [] args) {  
        findOysters();  
        findChampagne();  
        payAtCheckout();  
    }  
}
```

By contrast, our concurrent approach is expressed as:

```
class EfficientApproach {  
    class OysterFinder extends Thread {  
        public void run() {  
            findOysters(); rendezvousAtCheckout();  
        }  
    }  
    public static void main(String [] args) {  
        OysterFinder of = new OysterFinder();  
        of.start();  
        findChampagne(); rendezvousAtCheckout();  
    }  
}
```



Terminology

- A _____ is a computer with one processing core...
 - With a time-sharing OS, concurrent processing results in _____ (aka *logical concurrency*), because the OS time-shares the single core among the processes/threads.
- A _____ is a computer with multiple cores...
 - Concurrent processing results in _____ (aka *true concurrency*), as the OS can simultaneously run different processes/threads on different cores.
 - In a _____ *multiprocessor*, the cores
 - _____, and
 - are usually in *close physical proximity*.
 - In a _____ *multiprocessor*, the cores
 - have _____ (each has its own *local memory*),
 - are often *not in close physical proximity*.




Parallelizing Compilers

What *dependencies* exist in the following function?

```
void f(double x) {  
    double y, z;    // 1  
    fin >> y;       // 2 ... depends on 1  
    z = x * x;       // 3 ... depends on 1  
    return y * z;    // 4 ... depends on 2, 3  
}
```

Parallelizing compilers build _____,
and use them to identify pieces of code that can be
executed concurrently:

- branches in the graph (e.g., 2 & 3)
can be safely executed in parallel
- CPU must have the extra hardware to
perform _____



Processes and Threads

A computation is sometimes called a _____:

```
int main() {  
    s1;  
    s2;  
    ...  
    sn;  
}
```

The sequence of *events* that occur as control flows through a process is called a _____:

```
s1;  
s2;  
...  
sn;
```



If, for the same inputs, the sequence of events always occurs in the same order, the sequence is called _____; otherwise, the sequence is called _____.



Goal: Speedup

One goal of concurrent processing is to achieve speedup:

If a task requires τ time-units to solve sequentially,
but can be split into p subtasks that can be solved in parallel,
then parallel processing can perform it in _____ time-units.

Formally, speedup can be define as:

where: T_1 is the time 1 task takes to solve the problem, and
 T_N is the time N tasks take to solve the problem.

If it takes one person 10 minutes to find champaign+oysters,
but it takes two people 6 minutes, the speedup is _____.



Goal: Responsiveness

... is another goal of concurrent processing.

Example 1: _____

- If an application has a single thread and it has to perform a time-consuming task, then the GUI will “freeze” while the thread is performing that task.
- If the application uses one thread to handle user-interface events and forks a separate thread to perform each task, then the GUI will remain responsive.

Example 2: _____

- If a server has a single thread and has to perform a time-consuming task, then the server will be unable to accept incoming requests while it is performing the task.
- If the server uses one thread to accept incoming requests and forks a new thread to handle each request, the server will accept and handle all requests it receives.




Event Interleaving

When a computation consists of two or more processes:

```
void p() {
  a;
  b;
  c;
}
```










```
void q() {
  x;
  y;
}
```



then the events in the threads of those processes may be _____:

– Each sequence is _____,

but the *sequencing across both processes* is _____.



Interleaving in Ada

We can see non-deterministic behavior in Ada as follows:

```
procedure Interleave is
  task A;
  task body A is begin
    put("a");
  end A;
  task B;
  task body B is begin
    put("b");
  end B;
  task C;
  task body C is begin
    put("c");
  end C;
begin
  null;
end Interleave;
```

Sample Run:

```
% ./interleave
_____
% ./interleave
_____
% ./interleave
_____
% ./interleave
_____
...
```

How many distinct executions are there?

→ The number of _____ of {a, b, c}

→ Discrete Math: a set of n elements has _____ permutations, so _____ distinct possibilities.



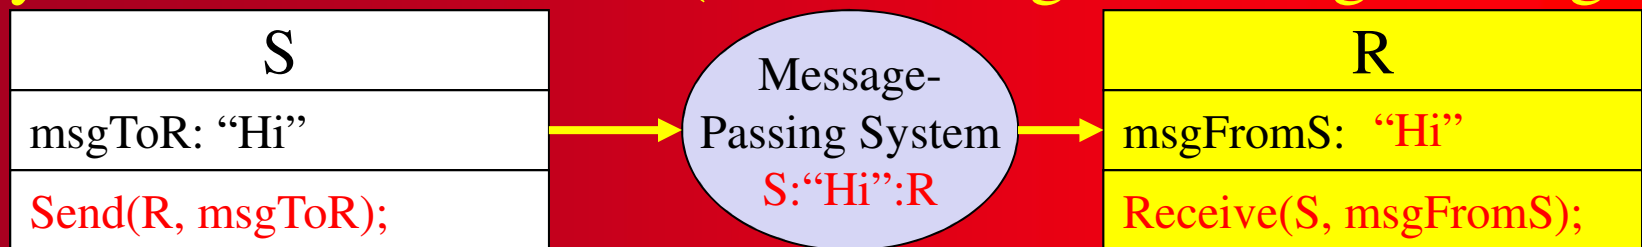
Communication

If a language allows us to divide a task into subtasks, it should also provide a way for those tasks to *communicate*.

- On a _____ multiprocessor, two tasks can communicate through the _____:



- On a _____ multiprocessor, they can communicate by _____ (i.e., sending-receiving messages):



Shared Memory Synchronization

If two tightly-coupled processes try to access shared memory simultaneously, the result may be incorrect... What can go wrong?

TravelAgent1		TravelAgent2
<pre>if (emptySeats > 0) emptySeats--; else display("no more seats!");</pre>	emptySeats: 1	<pre>if (emptySeats > 0) emptySeats--; else display("no more seats!");</pre>

Accesses to shared memory must be _____ to avoid this:

<pre>synchronize { // Java if (emptySeats > 0) emptySeats--; else display("no more seats!"); }</pre>	emptySeats: 1	<pre>synchronize { // Java if (emptySeats > 0) emptySeats--; else display("no more seats!"); }</pre>
--	---------------	--

Synchronization forces one process to _____ until the other is finished.



Synchrony


Concurrent computations lie on a continuum, depending on how much communication/synchronization they entail:

Problems whose solutions
require _____ synchrony

Problems whose solutions
require _____ synchrony

benefit of a concurrent solution

- *Lock-step synchronous* computations must communicate or be re-synchronized _____ of the computation:



- _____ (aka *embarrassingly parallel*) computations require no communication/synchronization of their processes




The _____ Problem

There are a number of “classic” synchronization problems, one of which is the producer-consumer problem...

There are M *producers* that put items into a buffer. The buffer is shared with N *consumers* that remove items from the buffer.

–The problem is to devise a solution that ensures _____.



Accesses to the buffer must be _____: if multiple producers / consumers access it simultaneously, producers may overwrite each other's values, consumers may retrieve the same value, etc.


In the *bounded-buffer version*, the buffer has a _____.



The _____ Problem

... is another “classic” synchronization problem:

Five philosophers sit at a table, alternating between eating noodles and thinking. In order to eat, a philosopher must have two chopsticks. However, there is a single chopstick between each pair of plates, so if one is eating, neither neighbor can eat. A philosopher puts down both chopsticks when thinking.



–Devise a solution that ensures:

- no philosopher starves; and
- a hungry philosopher is only prevented from eating by his immediate neighbor(s).


```
task philosopher is
  while True begin
    think(randomTime);
    get(left);
    get(right);
    eat();
    release(left);
    release(right);
  end while;
End philosopher;
```



The Dining Philosophers (ii)

How about this instead?

```
task philosopher is begin
  while True begin
    think(randomTime);
    while not (have(left) and have(right))
      begin
        get(left);
        if notInUse(right) then
          get(right);
        else
          release(left);
        end if;
      end while;
    eat;
    release(left);
    release(right);
  end while;
end philosopher;
```



Mutual Exclusion

An object that can only be accessed by “one-thing-at-a-time” (e.g., shared memory) is called a _____ object:

- An access by one process *excludes* other processes from access
- A task may have to ‘wait its turn’ to access the object
- Many real-world objects (e.g., chopsticks) are mutually exclusive
- Shared-memory *writes* are mutually exclusive: _____!

TravelAgent1		TravelAgent2
if (emptySeats > 0) emptySeats--; else display(“no more seats!”); end if;	emptySeats: 1	if (emptySeats > 0) emptySeats--; else display(“no more seats!”); end if;

- Shared-memory *reads* are not mutually exclusive, unless any task tries to write to the shared memory: _____!



Synchronization Primitives

In 1965, Dijkstra proposed the _____: a shared-memory programming mechanism that can be used to synchronize accesses to a mutually-exclusive resource, with two values: *{locked, unlocked}*, and three simple operations:

- *Initialize* the semaphore to *unlocked*
- *P*: *Lock* the semaphore (wait if it is already *locked*)
- *V*: *Unlock* the semaphore (awaken the first process waiting for it).

Java 1.7 added a *Semaphore* class:

Operation	Dijkstra	Java Syntax
Initialization (unlocked)	s: Semaphore ;	s = new Semaphore (1);
Lock the semaphore		
Unlock the semaphore		



Semaphores and Mutual Exclusion

A semaphore is a shared-memory variable that can be used to _____
_____ to other shared-memory variables:

TravelAgent1		TravelAgent2
<pre>_____ if (emptySeats > 0) emptySeats--; else display("no more seats!"); _____</pre>	<pre>emptySeats: 1 s: unlocked</pre>	<pre>_____ if (emptySeats > 0) emptySeats--; else display("no more seats!"); _____</pre>

Whichever travel agent executes $P(s)$ *first* (even by a nanosecond) will lock s , decrement $emptySeats$, and then unlock s .

Whichever travel agent executes $P(s)$ *second* will find s locked and have to wait until the other agent unlocks s , discover that $emptySeats == 0$, and then get the “no more seats” message.



Semaphores and Lockstep Synchrony

A semaphore also permits two processes to execute in lock-step:

S		R
msgToR:	sharedSpace:	msgFromS:
loop { _____ sharedSpace = msgToR; _____ }	okToWrite: unlocked okToRead: locked	loop { _____ msgFromS = sharedSpace; _____ }

- If *R* executes first, it will wait on the (*locked*) *okToRead* semaphore, until *S* signals (after it writes a value to the shared memory)...
- If *S* executes first, it will lock the (*unlocked*) *okToWrite* semaphore, write its message to shared memory, signal *okToRead*, and then wait on the (*locked*) *okToWrite* until *R* signals (after it finishes reading).


A _____ is a group of statements that access shared space.



Dining Philosophers (iii)

What if we use a Semaphore?

```
Semaphore: wantBothSticks;  
task philosopher is begin  
  while True begin  
    think(randomTime);  
    while not haveBothSticks begin  
      P(wantBothSticks);  
      if available(leftStick) and  
        available(rightStick) then begin  
        get(left);  
        get(right);  
      end if;  
      V(wantBothSticks);  
    end while;  
    eat();  
    release(leftStick);  
    release(rightStick);  
  end while;  
end philosopher;
```



This seems to do it, but synchrony is so tricky, it's hard to be 100% certain it's correct

...



Locks and Condition Variables

A semaphore may be used for either of two purposes:

- _____: guarding access to a critical section
- _____: making threads/processes suspend/resume

This dual use can lead to confusion: it may be unclear which role a semaphore is playing in a given computation...

For this reason, newer languages provide *distinct constructs for each*:

- _____: guarding access to a critical section
- _____: making threads wait until a condition is true

Locks support mutually-exclusive access to shared memory;
condition variables support thread/process synchronization.



Locks

Like a Semaphore, a lock has two associated operations:

- *acquire()* - try to lock the lock; if it is locked, go to sleep
- *release()* - unlock the lock; awaken a waiting thread (if any)

The *acquire()* is analogous to the Semaphore _____ operation;
the *release()* is analogous to the Semaphore _____ operation.

These can be used to 'guard' a critical section:

```
sharedLock.acquire();  
// access sharedObj  
sharedLock.release();
```

```
Lock sharedLock;  
Object sharedObj;
```

```
sharedLock.acquire();  
// access sharedObj  
sharedLock.release();
```

Every Java class inherits a *hidden lock* from class *Object*;
the _____ keyword uses it:

```
synchronized {  
    // critical section  
}
```



Condition Variables

A *Condition* is a predefined type available in some languages that can be used to declare variables for synchronization.

When a thread needs to suspend execution inside a critical section until some condition is met, a *Condition* can be used.

There are three operations for a *Condition*:

- _____
 - suspend immediately; enter a queue of waiting threads
- _____, aka *notify()* in Java
 - awaken a waiting thread (usually the first in the queue), if any
- _____, aka *notifyAll()* in Java
 - awaken all waiting threads, if any

Every Java class inherits it from class *Object* a hidden condition-variable, and the *wait()*, *notify()* & *notifyAll()* methods that use it.



Monitors

Semaphores, Locks, and Conditions are simple but powerful synchronization tools; but many believe that they are *too powerful for the average programmer* (like the *goto*)...

– _____ are easy mistakes to make

Just as control structures were “higher level” than the *goto*, language designers began looking for higher level ways to synchronize processes.

In 1973, Brinch-Hansen and Hoare proposed the _____, a class whose methods are automatically accessed in a mutually-exclusive manner.

– A monitor *prevents simultaneous access by multiple threads*



Mesa-Style Monitors

Concurrent Pascal was first to provide a *Monitor* construct:

```
type BoundedBuffer = Monitor
  constant N := 1024;
  myHead, myTail, mySize: integer := 0;
  myValues : array(0..N-1) of Object;
  notEmpty, notFull: Condition;

  procedure put(obj: Object) begin
    while mySize = N do notFull.wait; end;
    myValues(myHead) := obj;
    myHead := (myHead + 1) mod N; mySize := mySize + 1;
    notEmpty.signal;
  end;

  procedure get(var obj: Object) begin
    while mySize = 0 do notEmpty.wait; end;
    obj = myValues(myTail);
    myTail = (myTail + 1) % N; mySize := mySize - 1;
    notFull.signal;
  end;
end;
```




Monitor Visualisation

The compiler 'wraps' calls to *put()* and *get()* as follows:

... *call to put or get*

If the lock is *locked*, the thread enters the entry queue




Each cond. variable has its own internal queue, in which sleeping threads wait to be signaled...

public (interface)

hidden

private



Java Monitors

Java classes with _____ are monitors...

Example: Let's build a self-synchronizing *BoundedBuffer* class:

```
public class BoundedBuffer extends Object {
    private int mySize, myMax,
               myHead, myTail;
    private Object [] myValues;
    public BoundedBuffer(int n) {
        myMax = n;
        mySize = myHead = myTail = 0;
        myValues = new Object[n];
    }
    public synchronized int size() { return mySize; }
    public int capacity() { return myMax; }
    public synchronized int isFull() { return mySize == myMax; }
    public synchronized int isEmpty() { return mySize == 0; }
    // ... continued on next page ...
}
```

*A **synchronized** method
acquires the class's lock
to guarantee "one-thread-
at-a-time" execution...*



Buffer Synchrony

```
// ... continued from previous page ...
public synchronized void put(Object obj) {
    while ( this.isFull() )
        try{ wait(); } catch(Exception e) {}
    myValues[myHead] = obj;
    myHead = (myHead + 1) % myMax;
    mySize++;
    notifyAll();
}
public synchronized Object get() {
    Object result;
    while ( this.isEmpty() )
        try{ wait(); } catch(Exception e) {}
    result = myValues[myTail];
    myTail = (myTail + 1) % myMax;
    mySize--;
    notifyAll();
    return result;
}
}
```

The wait()
operation
causes the
executing
thread to

The
notifyAll()
operation


all waiting
threads.



Bounded Buffer Producer-Consumer

We can then use our *BoundedBuffer* as follows:

```
// producer thread
for (;;) {
    // produce Item it;
    buf.put(it);
}
```



```
// consumer thread
for (;;) {
    buf.get(it);
    // consume Item it;
}
```

_____ : No synchronization needed in producer or consumer.

Recall: Every Java class inherits a hidden *lock* from *Object*...

- When a *synchronized* method is called, it tries to *acquire* the lock (waiting if it is already locked), and *releases* the lock on termination.

Every Java class inherits a hidden *condition variable* from *Object*...

- `wait()` suspends a thread on the condition; `notify()` awakens a thread waiting on the condition; `notifyAll()` awakens all waiting threads.



More Recent Java

In 2005, Java 1.5 added the package _____:

- A `ThreadPool` class and an `Executor` framework to make the management of groups of threads easier and more convenient
- Classes for thread-safe data structures (list, queue, map, ...)
- Classes for synchronization (semaphore, barrier, mutex, latch, ...)
- Classes for creating lock and condition variables;
- Classes for atomic operations (arithmetic, test-and-set, ...)

Subsequent Java releases have continued to add features:

- Futures, for asynchronous computations
- `ForkJoinTasks` and `ForkJoinPools` for recursive parallelism
- Lambda expressions, `CompletableFuture`, parallel streams, `WorkStealingThreadPools` for load-balancing, ...



OpenMP ...

- ... stands for *Open MultiProcessing*
- ... is an industry-standard library for shared-memory parallel computing in C, C++, Fortran, ...
- ... uses _____
- ... simplifies the task of parallelizing legacy code
- ... was designed by a large consortium in 1997:
AMD, Cray, Fujitsu, HP, IBM, Intel, NEC, Nvidia, Oracle, Redhat, TI, ...
- ... has “built in” support for many parallel design patterns
- ... continues to evolve (OpenMP 2.0 in 2000; 3.0 in 2008; 4.0 in 2014, ...; current version is 4.5)



Example: Summing an Array

```
#include <iostream>
#include <omp.h>

// getFileName(), readFile(), ...

int main(int argc, char** argv) {
    string filename = getFileName(argc, argv);
    vector<int> v = readFile(filename);

    #pragma omp parallel for reduction(+:sum)
    for (int i = 0; i < v.size(); i++) {
        sum += v[i];
    }

    cout << "The sum of the values in '"
         << fileName << "' is " <<
         << sum << endl;
}
```

On a machine
with t cores,
this directive

_____ threads

An _____
_____ occurs
at the end of
the stmt that
follows the
#pragma




#pragma omp parallel for


The #pragma omp parallel directive forks the threads...

The for clause _____ that follows it across those threads.


On a single-core machine: _____ thread...



On a dual-core machine: _____ threads...

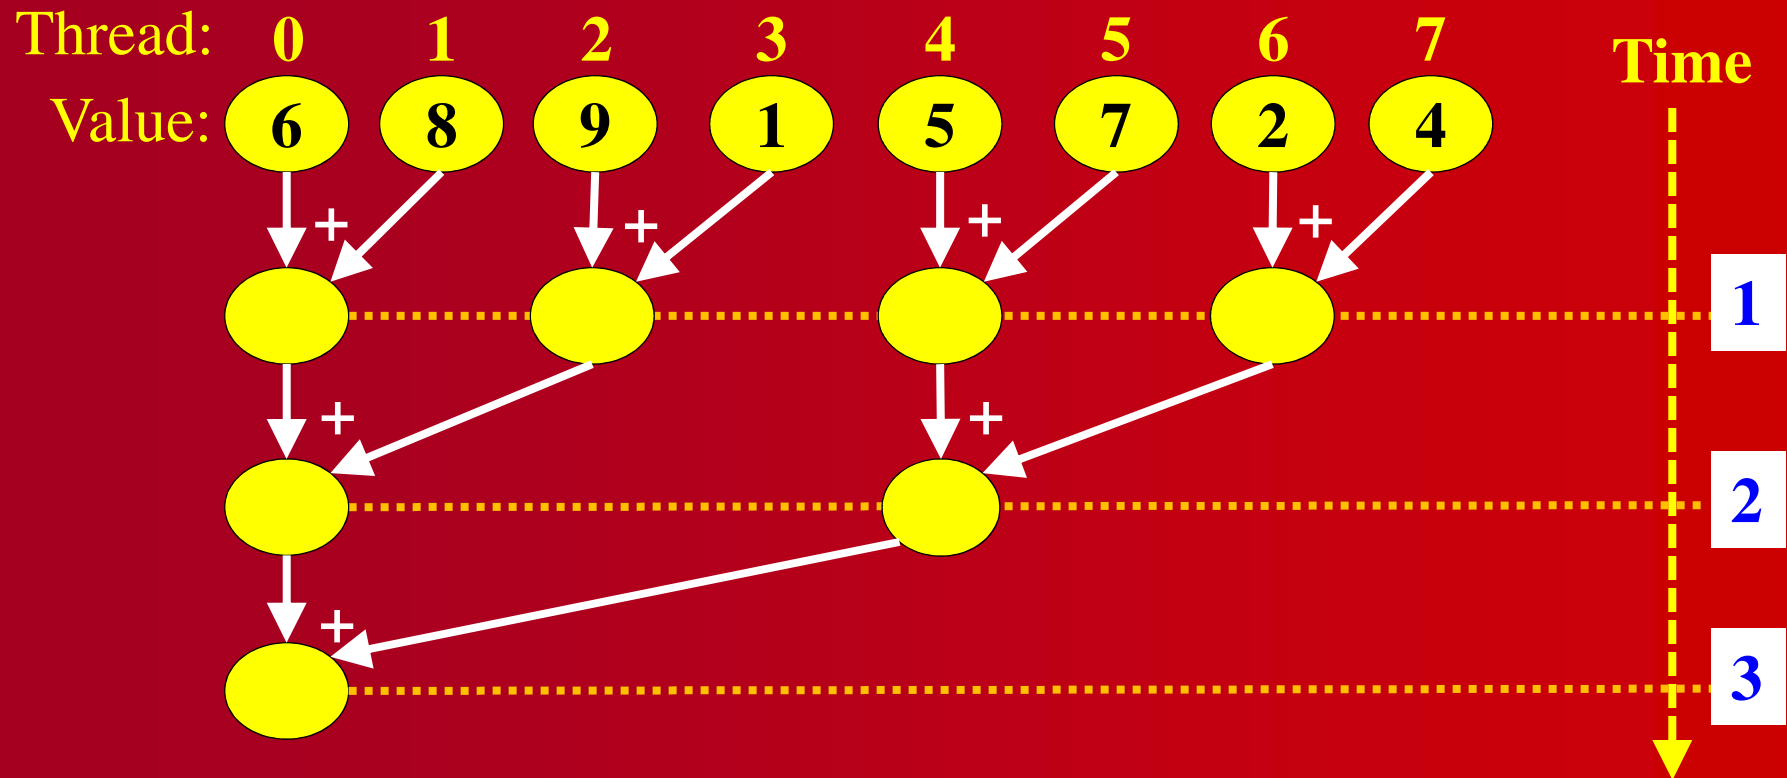


On a quad-core machine: _____ threads...



The Reduction Clause

... in the `#pragma omp parallel for reduction(+:sum)` uses the `+` operator to combine the partial sums into variable `sum`.



Reduction reduces the sum-time from $O(t)$ to _____ ...



OpenMP also supports...

Other ways to reduce local results: +, *, -, /, &, |, ^, &&, ||, ...

Other directives: *#pragma omp* _____

- *critical*
- *sections*
- *task*
- *single*
- *atomic*
- *section*
- *teams*
- *master*
- *barrier*
- *taskwait*
- *simd*
- ...

Library functions:

- *omp_set_num_threads()*
- *omp_get_num_threads()*
- *omp_get_thread_num()*
- *omp_get_num_procs()*
- *omp_init_lock()*
- *omp_set_lock()*
- *omp_unset_lock()*
- *omp_test_lock()*
- *omp_get_num_teams()*
- *omp_get_team_size()*
- *omp_get_wtime()*
- ...

Much, much more!

